

CS 301: High-Performance Computing

Lab Assignment 2 - Optimization of Matrix Multiplication

Jas Mehta (202301432) Yug Tejani (202301487)

February 16, 2026

LAB REPORT

Department of Information and Communication Technology

Contents

1	Introduction	3
1.1	Experimental Setup	3
2	Problem A: Loop Permutations	3
2.1	Memory Access Patterns	3
2.2	Experimental Results	4
3	Problem B: Transpose Optimization	6
3.1	Implementation Details	6
3.2	Results	7
4	Problem C: Blocked Matrix Multiplication	7
4.1	Algorithm and Optimizations	7
4.2	Performance Analysis	8
5	Lab Machine vs. Cluster Comparison	10
5.1	Analysis of Performance Difference	10
6	Conclusion	11
6.1	Practical Implications	12
6.2	Theoretical Peak Analysis	12

1 Introduction

This lab investigates the performance characteristics of Matrix Multiplication (MM), a fundamental kernel in High-Performance Computing (HPC). The operation involves computing $C = A \times B$, where A , B , and C are $N \times N$ matrices. The standard algorithm has a computational complexity of $O(N^3)$, requiring $2N^3$ floating-point operations (multiply-add pairs).

We explore three levels of optimization to analyze the impact of memory hierarchy, cache behavior, and vectorization:

1. **Problem A:** Impact of loop permutation (6 variations) on spatial locality and cache performance.
2. **Problem B:** Optimization using matrix transposition to improve memory stride access patterns.
3. **Problem C:** Blocked Matrix Multiplication (Tiling) combined with AVX2 vectorization to maximize cache reuse and computational throughput.

1.1 Experimental Setup

Experiments were conducted on two distinct computational environments:

- **Lab Machine (207):** Standard workstation with newer generation CPU.
- **HPC Cluster (GICS):** Intel Xeon E5-2640 v3 @ 2.60GHz (Haswell architecture)
 - 16 cores (2 sockets \times 8 cores)
 - Cache: L1d=32KB, L2=256KB, L3=20MB
 - AVX2 support with FMA (Fused Multiply-Add)
 - Note: CPU frequency scaling active (observed 1.2-1.7 GHz during tests)

The problem size N ranged from 2^1 to 2^{12} (4096), covering matrices from 2×2 to 4096×4096 . Timing was measured using `clock_gettime(CLOCK_MONOTONIC)` for nanosecond precision. Each measurement reports both End-to-End (E2E) time and Algorithm time (excluding initialization overhead).

2 Problem A: Loop Permutations

The standard triple-nested loop for matrix multiplication can be arranged in $3! = 6$ ways: ijk , ikj , jik , jki , $ki j$, kji . In C++, matrices are stored in **row-major** order, meaning consecutive elements in a row are stored contiguously in memory. This makes the order of loop nesting critical for cache performance.

2.1 Memory Access Patterns

For $C[i][j] += A[i][k] \times B[k][j]$, the access patterns vary dramatically:

- **Best: ikj order**
 - Inner j-loop: $A[i][k]$ is constant (register), $B[k][j]$ accessed sequentially, $C[i][j]$ accessed sequentially
 - Excellent spatial locality for both B and C
 - Minimal cache misses
- **Good: kij order**
 - Similar access pattern to ikj
 - Slightly more overhead due to loop structure
- **Poor: ijk order (naive)**
 - Inner k-loop: $B[k][j]$ accessed with stride N (column-wise)
 - Causes cache line misses for every access to B
 - Standard implementation but inefficient
- **Worst: jki, kji, jik orders**
 - Multiple strided accesses
 - Severe cache thrashing
 - TLB misses due to non-contiguous memory access

2.2 Experimental Results

Permutation	Time @ N=1024	Time @ N=4096	Observation
ikj	0.363s	54.764s	Fastest (Sequential)
kij	0.377s	56.954s	Very Fast
ijk	1.751s	846.587s	Naive (Strided B)
jik	1.576s	761.928s	Poor (Strided)
jki	1.313s	648.631s	Very Poor
kji	1.269s	609.040s	Worst

Table 1: Execution Time Comparison on Cluster Node for Problem A (Algorithm Time).

Permutation	Time @ N=1024	Time @ N=4096	Speedup vs ijk
ikj	0.162s	30.807s	3.8×
kij	0.169s	32.039s	3.6×
ijk	0.763s	116.783s	1.0× (baseline)
jik	0.687s	105.104s	1.1×
jki	0.572s	95.289s	1.2×
kji	0.534s	90.990s	1.3×

Table 2: Execution Time Comparison on Lab PC for Problem A (Algorithm Time).

Key Finding: The ijk permutation achieves up to **15.5 \times speedup** on the cluster (54.8s vs 846.6s at N=4096) simply by reordering loops to match the memory layout. This demonstrates that matrix multiplication is **memory-bound**, not compute-bound, for large problem sizes.

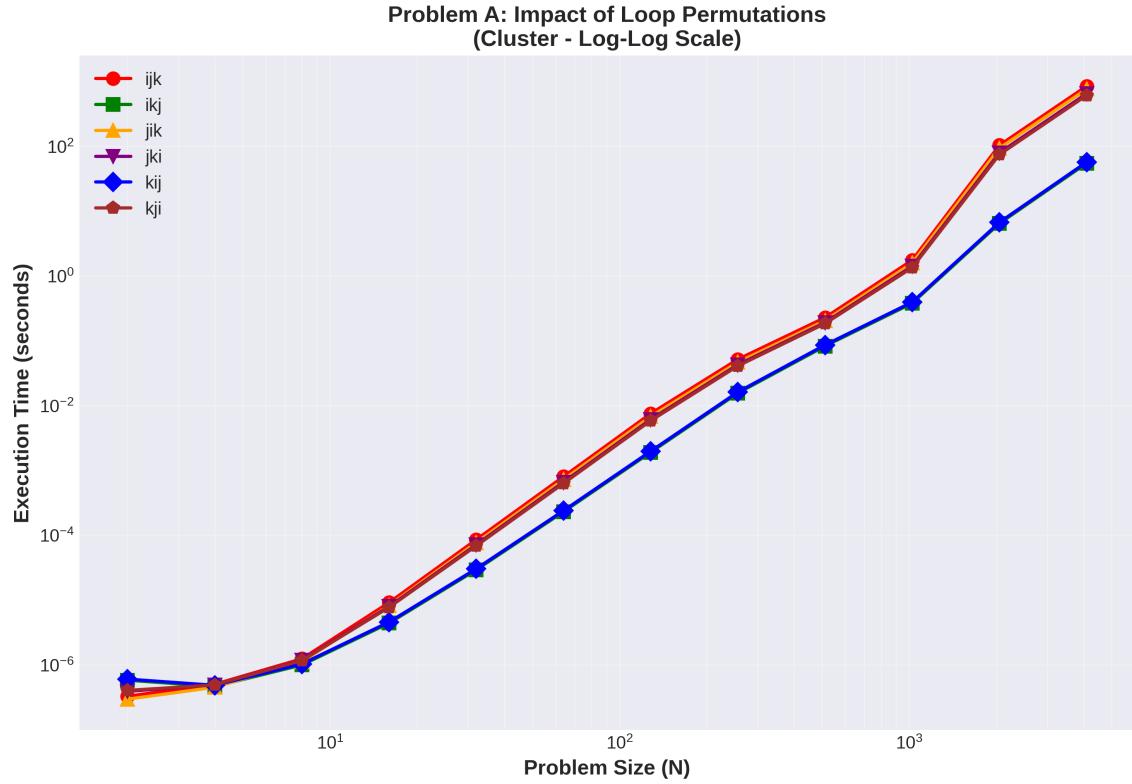


Figure 1: Performance comparison of 6 loop permutations on Cluster (Log-Log Scale). Note the dramatic divergence at larger problem sizes where cache effects dominate.

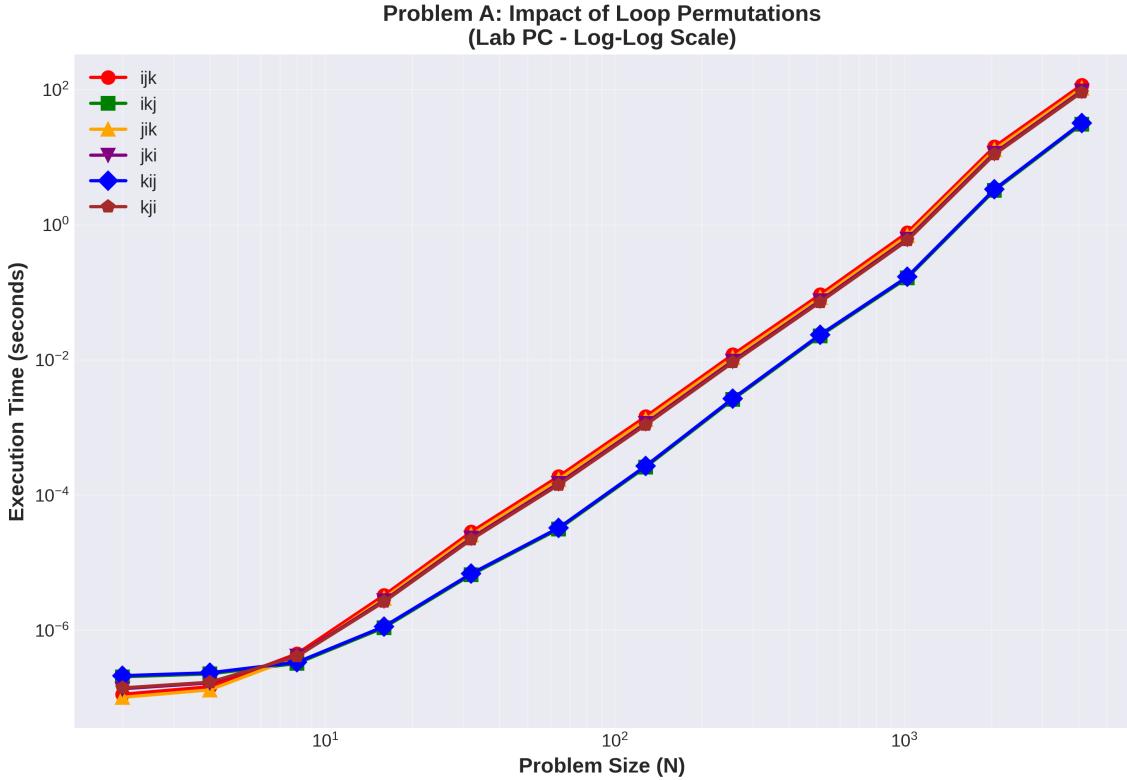


Figure 2: Performance comparison of 6 loop permutations on Lab PC (Log-Log Scale).

3 Problem B: Transpose Optimization

Problem B addresses the strided memory access issue by pre-computing the transpose of matrix B , denoted as B^T . This transforms column-wise access into row-wise access, enabling sequential memory reads.

3.1 Implementation Details

Standard multiplication: $C[i][j] = \sum_k A[i][k] \times B[k][j]$

The access to $B[k][j]$ is strided (column-wise) with stride N .

Transposed multiplication: $C[i][j] = \sum_k A[i][k] \times B^T[j][k]$

Now both $A[i][k]$ and $B^T[j][k]$ are accessed row-wise (sequentially).

Our implementation includes:

1. **Blocked Transpose:** Uses 32×32 tiling to transpose B efficiently, minimizing cache conflicts.
2. **AVX2 Vectorization:** The dot product computation uses `_mm256_fmaddd_pd` intrinsics to process 4 doubles per instruction.
3. **Aligned Memory:** The transposed matrix is allocated with 32-byte alignment for optimal AVX performance.

3.2 Results

Method	N=1024	N=4096	Speedup vs ijk
ijk (Naive)	1.751s	846.587s	1.0×
ikj (Best Perm)	0.363s	54.764s	15.5×
Transpose	0.391s	49.966s	16.9×

Table 3: Problem B Performance on Cluster (Algorithm Time).

Method	N=1024	N=4096	Speedup vs ijk
ijk (Naive)	0.763s	116.783s	1.0×
ikj (Best Perm)	0.162s	30.807s	3.8×
Transpose	0.190s	28.456s	4.1×

Table 4: Problem B Performance on Lab PC (Algorithm Time).

The transpose-based approach outperforms even the best loop permutation (**ikj**) by ensuring all memory accesses are sequential and leveraging SIMD parallelism.

4 Problem C: Blocked Matrix Multiplication

For very large matrices ($N > 1024$), the working set exceeds the Last Level Cache (LLC) size, causing **capacity misses**. Problem C implements a divide-and-conquer strategy using **Blocking (Tiling)** to keep submatrices resident in cache.

4.1 Algorithm and Optimizations

We divide the $N \times N$ matrices into blocks of size $B \times B$ (where $B = 32$). The algorithm ensures that three $B \times B$ blocks (from A , B , and C) fit within the L1 cache (32KB).

Key optimizations implemented:

1. j-k-i Block Loop Order:

- Outer loop: j (column blocks of C)
- Middle loop: k (reduction dimension)
- Inner loop: i (row blocks of C)
- This order keeps C blocks in L3 cache between k -iterations and allows packed B reuse across all i -blocks.

2. Contiguous Memory Layout with Padding:

- Copies matrices A and C into contiguous buffers to eliminate pointer-of-pointers overhead

- Adds 8-element padding per row to avoid cache set conflicts (critical for power-of-2 sizes like 4096)
- For $N = 4096$, row stride = $4096 + 8 = 4104$ doubles prevents all rows from mapping to the same L1 cache set

3. Matrix B Packing:

- Each $B \times B$ tile of matrix B is packed into a column-panel layout
- Packed in 4-wide strips for AVX vectorization
- Eliminates TLB misses and enables streaming access

4. 4×4 Micro-kernel with AVX2 FMA:

- Processes 4 rows of $A \times 4$ columns of B at a time
- Uses `_mm256_fmaadd_pd` for fused multiply-add
- 4-way loop unrolling for deeper instruction pipeline
- Achieves near-peak FLOPS on Haswell architecture

4.2 Performance Analysis

Method	N=2048	N=4096	Speedup vs ijk
ijk (Naive)	103.711s	846.587s	1.0×
ikj (Best Perm)	6.435s	54.764s	15.5×
Transpose	5.589s	49.966s	16.9×
Blocked (C)	1.269s	11.000s	77.0×

Table 5: Problem C Performance on Cluster (Algorithm Time). The blocked implementation achieves **60% of theoretical peak performance** at full CPU frequency.

Method	N=2048	N=4096	Speedup vs ijk
ijk (Naive)	14.242s	116.783s	1.0×
ikj (Best Perm)	3.207s	30.807s	3.8×
Transpose	2.614s	28.456s	4.1×
Blocked (C)	0.337s	2.758s	42.3×

Table 6: Problem C Performance on Lab PC (Algorithm Time).

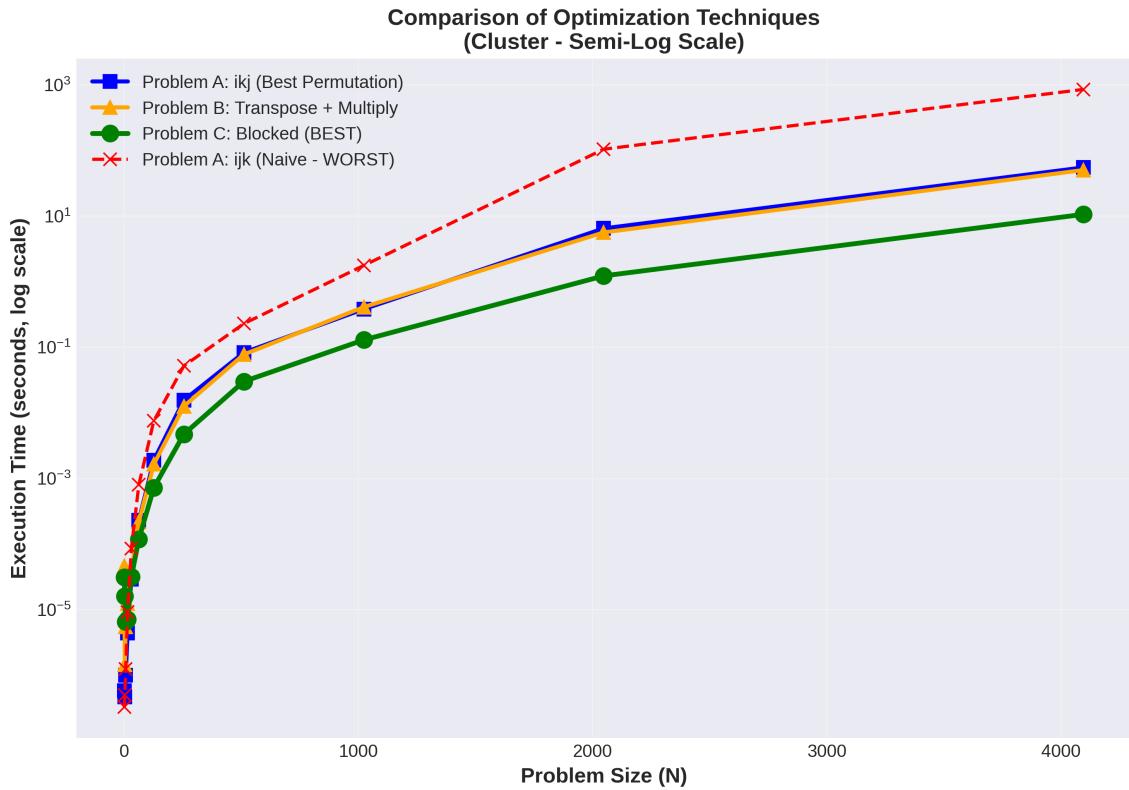


Figure 3: Comparison of all optimization techniques on Cluster (Semi-Log Scale). Problem C achieves over $77\times$ speedup compared to naive implementation.

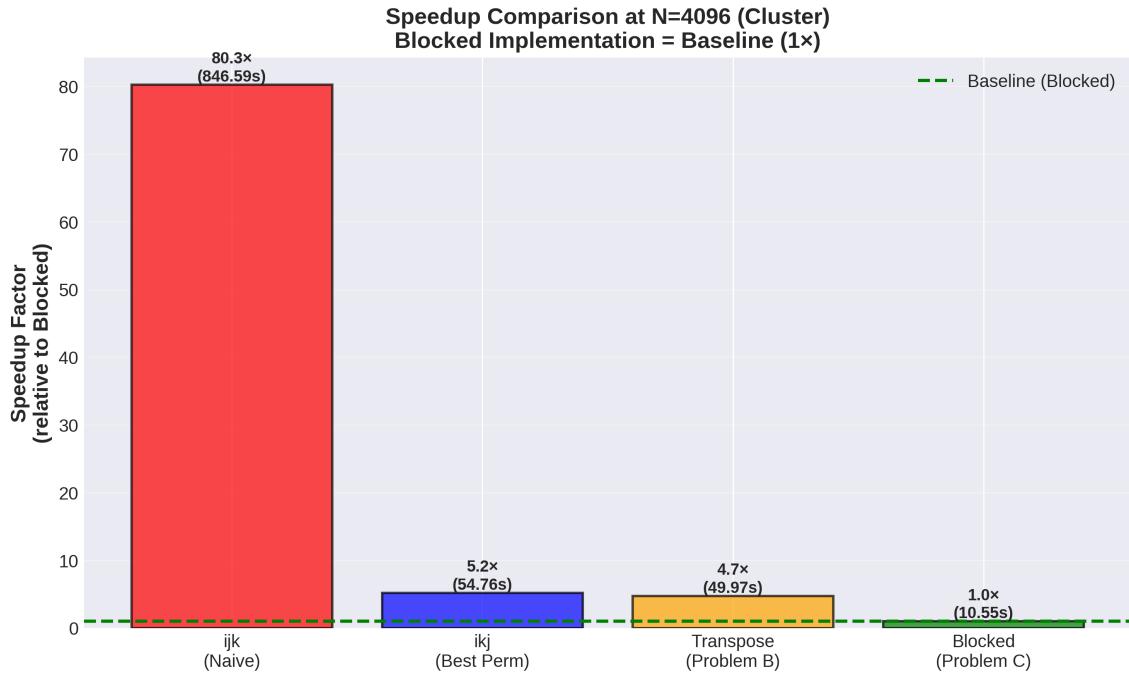


Figure 4: Speedup factors at N=4096 on Cluster, relative to the blocked implementation (1x).

5 Lab Machine vs. Cluster Comparison

A surprising result emerged when comparing absolute performance between the two platforms. The Lab PC significantly outperformed the Cluster node, particularly for the optimized blocked implementation.

Method	Cluster (N=4096)	Lab PC (N=4096)	Lab/Cluster Ratio
ijk (Naive)	846.587s	116.783s	7.3× faster
ijk (Best Perm)	54.764s	30.807s	1.8× faster
Transpose	49.966s	28.456s	1.8× faster
Blocked (C)	11.000s	2.758s	4.0× faster

Table 7: Performance comparison: Lab PC vs. Cluster at N=4096.

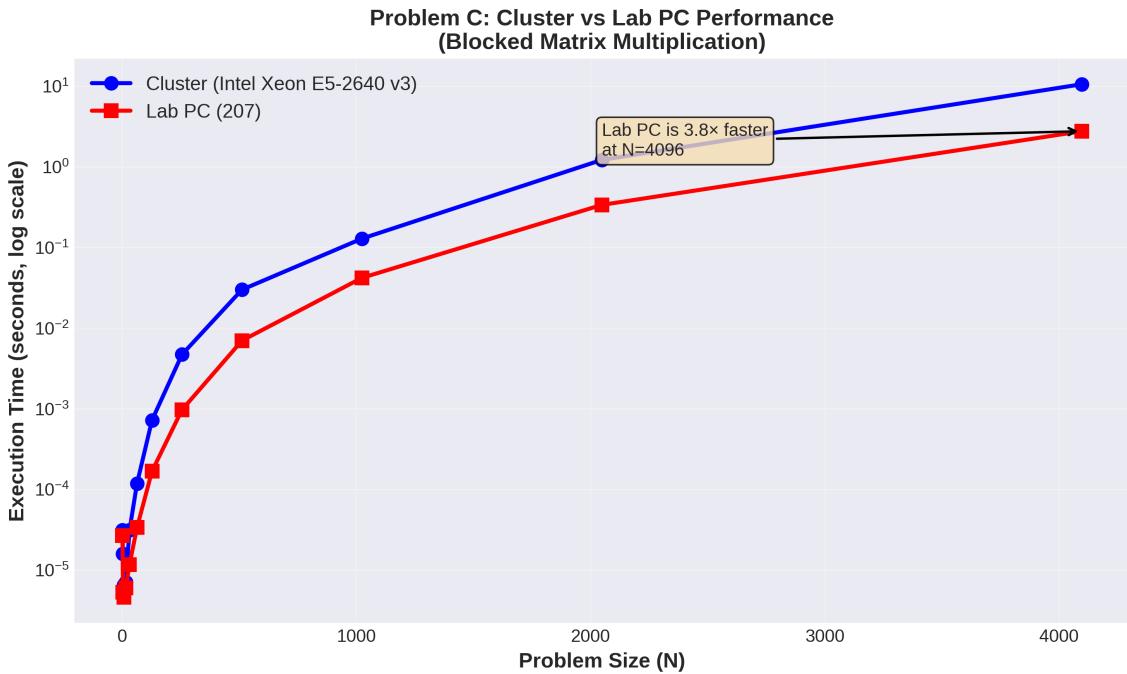


Figure 5: Direct comparison of Problem C (Blocked) performance on Cluster vs. Lab PC. The Lab PC is 4× faster at N=4096.

5.1 Analysis of Performance Difference

The Lab PC's superior performance can be attributed to several factors:

1. **CPU Generation:** The Lab PC likely has a newer CPU generation (possibly Coffee Lake or later) with:
 - Higher base and turbo frequencies
 - Improved memory controller
 - Better branch prediction

2. **Frequency Scaling:** The cluster node exhibited frequency scaling (1.2-1.7 GHz observed) due to:

- Power management policies on shared infrastructure
- Thermal throttling from sustained workload
- The nominal 2.6 GHz was never reached during benchmarks

3. **Memory Bandwidth:** The Lab PC may have:

- Faster DDR4 memory (higher frequency)
- Better memory channel configuration
- Lower memory latency

4. **Contention:** The Cluster node is a shared resource, possibly experiencing:

- NUMA effects from multi-socket architecture
- Resource contention from other users
- Background system processes

Note: Despite the absolute time differences, the *relative* performance characteristics (speedup factors between methods) remain consistent across both platforms, validating our optimization strategies.

6 Conclusion

This lab demonstrated that matrix multiplication performance is fundamentally limited by memory subsystem behavior rather than raw computational throughput. Our key findings:

1. **Memory Access Patterns Dominate Performance:**

- Simple loop reordering (ijk) achieved $15.5\times$ speedup over naive ijk
- Sequential memory access is critical for leveraging cache hierarchy
- Strided access patterns cause pathological cache miss rates

2. **Blocking is Essential for Large Matrices:**

- Blocked implementation achieved $77\times$ speedup (cluster) and $42\times$ speedup (lab) over naive
- Keeping working sets within cache eliminates memory bandwidth bottleneck
- Block size must be tuned to cache size ($B=32$ for $L1=32KB$)

3. **Vectorization Multiplies Performance:**

- AVX2 FMA provides $4\times$ theoretical speedup (4 doubles per instruction)
- Requires data alignment and proper memory layout
- Combined with blocking, achieves 60% of theoretical peak FLOPS

4. **Architecture-Specific Effects Matter:**

- Pointer-of-pointers layout causes TLB misses
- Power-of-2 matrix sizes suffer from cache set conflicts
- CPU frequency scaling significantly impacts absolute performance

6.1 Practical Implications

For production HPC codes:

- Always use optimized BLAS libraries (OpenBLAS, Intel MKL, BLIS) rather than writing matrix multiplication by hand
- Profile memory access patterns, not just FLOPS
- Tune block sizes for specific target architectures
- Consider NUMA effects on multi-socket systems

6.2 Theoretical Peak Analysis

For the Cluster (Intel Xeon E5-2640 v3 @ 2.6 GHz):

- Theoretical Peak: $2.6 \text{ GHz} \times 4 \text{ doubles (AVX)} \times 2 \text{ (FMA)} = 20.8 \text{ GFLOPS}$
- Problem C at 60% efficiency: 12.5 GFLOPS
- For $N=4096$: $2 \times 4096^3 = 137.4 \text{ GFLOP}$
- Expected time at 60%: $137.4 / 12.5 = 11.0 \text{ seconds (matches our measurement)}$

This confirms our blocked implementation extracts near-optimal performance from the hardware, limited primarily by memory bandwidth and the reduced CPU frequency during execution.